satishfulwani.in
github.com/satishf889
satish.fulwani63@gmail.com
+91-9145175111

# SATISH FULWANI

Senior Data Engineer | Snowflake, Kafka, Azure & AWS | 3x AWS Certified
Pune, India · MTech Data Science & Engineering, BITS Pilani

## SUMMARY

**Staff Data Engineer** with **8+ years** building reliable, cloud-native data platforms. Seeking **Senior Data Engineer** roles at large investment firms. Strengths include **Snowflake** warehousing, **realtime pipelines** (Kafka, Lambda), batch ETL/ELT on **Azure Data Factory** and **Databricks**, and production **GenAI/RAG** on **Snowflake Cortex AI** and **LangChain**.

Led a team of **3** to modernize a legacy platform on **Azure and Snowflake**, migrating **1+ TB** with **zero downtime** for a major **regulated insurance** program. Delivered agent **skills and automations** and event-driven streaming into Snowflake for analytics. Mentor engineers; partner with product and business on data quality and SLAs.

Proven delivery in regulated, high-volume environments (insurance, media, nutrition); comfortable partnering with business and compliance stakeholders on **data governance**, quality SLAs, and platform reliability.

Rapidly adapts to new stacks, tooling, and domains - learn, unlearn, and relearn is a working principle.

## TECHNICAL SKILLS

**DATA & STREAMING:**
**Snowflake**, **Kafka**, PySpark, Databricks, SQL, Azure Data Factory, Pandas, data modeling

**CLOUD:**
AWS (Lambda, Glue, S3, API Gateway, DynamoDB), Azure (ADF, Databricks, Synapse, Container Apps), **Terraform**, **Kubernetes**, Jenkins CI/CD, Serverless

**GENAI & TOOLING:**
Python, **LangChain**, RAG, Snowflake Cortex AI, vector search, agent skills & automations, data governance & lineage

## AWARDS & RECOGNITION

- **Brightest Mind** - Nagarro's company-wide engineering excellence award.
- 2nd place, Accenture AI Hackathon (Aug 2019) - AI solution pitched and demoed to senior leadership.
- **APEX Award** x2 - Accenture's highest performance award, for Delivery and for Margin.
- Built HR mail automation system in Python, adopted company-wide by the HR team.

## CERTIFICATIONS

- GCP Certified Cloud Digital Leader | GCP | Dec'22
- AWS Certified Machine Learning - Specialty | AWS | Mar'20
- AWS Certified Cloud Practitioner | AWS | Mar'20
- AWS Certified Developer - Associate | AWS | Sep'19
- Machine Learning Advanced Nanodegree | Udacity | Jan'19
- Machine Learning Foundation Nanodegree | Udacity | Apr'18

## EDUCATION

**BITS PILANI**
MTECH DATA SCIENCE & ENGINEERING | APR'21-APR'23 | PILANI, RJ
CGPI - 7.18/10

**MUMBAI UNIVERSITY**
B.E. COMPUTER SCIENCE | JUN'14-JUN'18 | VESIT, MUMBAI, MH
CGPI - 7.19/10

## PROFESSIONAL EXPERIENCE

### Nagarro India Pvt Ltd

Feb'23 - Present | Pune, IN

#### Role: STAFF DATA ENGINEER

**Industry:** Large Consultancy Firm
**Project:** AWS Data Platform, GenAI Bot, Realtime Streaming & Agent Automations
**Technologies:** AWS, Terraform, LangChain, Kafka, Lambda, Snowflake Cortex AI, S3, API Gateway, Python

Led architecture and delivery of a Terraform-managed AWS data platform, a production GenAI bot, and realtime Kafka streaming. Mentored engineers and drove adoption of agent-tooling patterns across the team.

**Key Achievements:**

- Designed and provisioned end-to-end AWS infrastructure with **Terraform** (Lambda, S3, API Gateway, IAM policies), ensuring reproducible, secure deployments.
- Architected and deployed a production **RAG chatbot** on **Snowflake Cortex AI**; engineered **LangChain**-driven document parsing, chunking, and embedding into **Snowflake native vector tables** for semantic retrieval.
- Built a **Kafka** event-ingestion layer with **Lambda** consumer workers curating near-real-time datasets in **Snowflake** for analytics and downstream BI dashboards.
- Designed and published reusable **agent skills and automations** covering code review, documentation, and operational runbooks - measurably reducing team delivery cycle time.
- Mentored engineers on GenAI patterns and RAG design; established team-wide best practices for agent tooling and observability.

#### Role: ASSOCIATE STAFF DATA ENGINEER

**Industry:** Asia's Largest Insurance Company
**Project:** Azure Cloud Data Pipeline Modernization
**Technologies:** Azure Data Factory, Databricks, PySpark, Azure SQL, Snowflake

Led a team of 3 to modernize Asia's largest insurance company's on-premises data estate onto Azure, delivering reliable pipelines, regulatory-grade governance, and executive analytics.

**Key Achievements:**

- Led a **team of 3** to migrate **1+ TB** of legacy insurance data to **Azure SQL** and **Snowflake** with **zero downtime**, coordinating cutovers across business and ops stakeholders.
- Designed and delivered **Azure Data Factory** orchestration pipelines and **Databricks** / **PySpark** ETL workflows processing large-scale, regulated insurance data at production SLAs.
- Implemented data quality frameworks, pipeline observability, and lineage tracking across **20+ ADF pipelines** to meet regulatory compliance and audit requirements.
- Partnered with solution architects and business stakeholders to translate complex insurance-domain requirements into reliable, versioned **data contracts** and **SLA-bound** delivery.
- Enforced **data governance** policies and access controls across Azure and Snowflake workloads; reduced pipeline incident response time through alerting and runbook automation.

### Deloitte ILLP

Aug'21 - Feb'23 | Mumbai, IN

#### Role: DATA ENGINEER CONSULTANT

**Industry:** Dutch Nutrition Company
**Project:** Solving Complex Big Data Issues Based on Business Needs
**Technologies:** Spark (PySpark), Python, AWS Glue, Hadoop, AWS S3, Microsoft Dynamics (CRM)

Aligned data engineering with evolving business needs: data quality, optimized Glue/Athena/S3 lakehouse patterns, and SageMaker notebooks for ML workflows serving a global nutrition client.

**Key Achievements:**

- Architected and delivered ETL on **AWS Data Lake** (**Glue**, Crawlers, **Athena**) and ML-bound pipelines feeding multiple production models.
- Delivered near-real-time executive dashboards; enabled **SageMaker** ML workflows on lakehouse data.

### Accenture Pvt Ltd

Oct'18 - Aug'21 | Bangalore, IN

#### Role: BIG DATA ENGINEER

**Industry:** US based Media & Entertainment Client
**Project:** Implemented migration of static files to AWS S3, optimizing storage and delivery efficiency
**Technologies:** AWS Lambda, SQS, DynamoDB, S3, EC2

Designed economical AWS migration architecture for on-premises media ingestion using serverless, fault-tolerant workflows while preserving data integrity for a US media and entertainment client.

**Key Achievements:**

- Architected event-driven pipelines with **Lambda**, **SQS**, and **DynamoDB** for resilient, scalable ingestion.
- Migrated **1.5 million** media assets to **S3** with ~**$100** migration overhead via serverless design.

## SELECTED PROJECTS

### Personal Projects

#### AI Travel Planning Assistant

LangChain · Azure OpenAI · ChromaDB · Streamlit · Docker

- **TravelAgent** orchestrator: **RAG** (**ChromaDB**) + **MCP** tools; **Streamlit**, **Docker Compose**, **pytest**.

#### Event Management Platform

AI-assisted development

- Event platform with **AI chatbots**, connections, attendance, and networking on **Azure Container Apps** / **Kubernetes**; end-to-end **AI-assisted development**.

#### Fruit Classifier (MTech)

Deep Learning · AWS/Keras | Apr'23

- MTech thesis: CNN fruit classification via **transfer learning** (ResNet/VGG) on **AWS**/**Keras**; compared accuracy and cost across architectures.